

Slide Designer — Three-Week Delivery Plan

What We're Building

Over the next three weeks the team is evolving the Slide Designer from a single-slide prototype into a production-quality tool. By the end of Week 3 the system will be able to process entire decks in parallel, handle all major slide types accurately (tables, images, diagrams, text layouts), and include an independent quality-checking module that can evaluate any slide — even outside the designer pipeline.

Team

Person	Role
Rafał	Team Lead — AI quality, prompt engineering, architecture
Max	Rendering engine — slide assembly, diagrams, text rendering
Harsh	Image handling
Shaz	Quality judge module
Availability note: Rafał will be unavailable during Week 2. All AI-dependent work (prompt tuning, agent logic) is front-loaded into Week 1 and back-loaded into Week 3. Week 2 is structured so Max, Harsh, and Shaz can work fully independently.	

Week 1 — Foundation

Goal: Fix the most impactful quality issues, improve table output, and lay the groundwork for processing full decks.

What each person is working on

Rafał	- Fixes several known accuracy bugs that are causing incorrect output on non-standard slides (already completed before this plan started) - Improves how the AI plans and renders text-heavy slides — better spacing, bullet formatting, decorative elements - Ports the dedicated table-processing logic from the previous version of the tool into the new pipeline (in progress)
Max	Builds the infrastructure to process multiple slides from a deck in one command, running them in parallel
Shaz	- Sets up the standalone quality judge module - Collects 20+ before/after slide pairs to use as a test dataset for measuring judge accuracy

What to expect at the end of Week 1

- Single-slide output is noticeably better: text containers look more professional, tables have correct structure
- The tool accepts a list of slides as input (e.g. “process slides 1, 3, and 5–8 from this deck”)
- The quality judge module exists as a standalone project, ready for Shaz to iterate on independently

Week 2 — Quality Features

Goal: Add image handling, improve text rendering, and add support for block diagrams. Rafał is absent this week — all work is

self-contained and does not require AI prompt changes.

What each person is working on

Max

- Improves how bullet points and mixed text styles (bold, italic, colour) are rendered in the output
- Adds support for block diagrams: boxes connected by arrows, which currently do not render correctly

Harsh

- Implements image passthrough: images embedded in the source slide are now carried through to the redesigned output (currently they are lost)

Shaz

- Runs the judge against the test dataset to establish a baseline accuracy score
- Iterates on the judge's rules and prompts to reduce false approvals and false rejections
- Target: judge correctly classifies at least 80% of test cases

What to expect at the end of Week 2

- Slides with images produce output that includes those images in the correct position
- Text slides with bullets and mixed formatting look correct
- Block diagram slides render with visible nodes and connecting arrows
- The quality judge reaches $\geq 80\%$ accuracy on the test dataset

Week 3 — Polish and Integration

Goal: Add mimic mode, finalise the judge, run end-to-end deck tests, and refine quality with professional slide makers. Rafał returns this week.

What each person is working on

Rafał

- Implements **mimic mode**: a new option that preserves the layout and structure of the source slide rather than redesigning it from scratch — useful when the source is nearly correct and only needs enhancement
- Adds AI prompt improvements for block diagrams (building on Max's rendering work from Week 2)
- Works with the external professional slide making company to review output quality and implement their feedback
- Integrates the finalised judge module into the main pipeline

Max

- Merges all branches and runs the full end-to-end deck pipeline test (10-slide deck, all features active)
- Final CLI polish: progress reporting, dry-run mode, multi-slide output directory

Harsh

- Supports accuracy refinement as needed

Shaz

- Finalises the judge module: configurable thresholds, final evaluation run, documentation
- Hands off a self-contained, documented module ready for independent use or integration

What to expect at the end of Week 3

- A full 10-slide deck can be processed end-to-end in under 5 minutes
- **Mimic mode** is available: the tool can enhance a slide while preserving its original layout
- All slide types work correctly: text, tables, images, block diagrams

- The quality judge is integrated and can be enabled as a pipeline gate or used as a standalone slide evaluation tool
- Professional slide makers rate the output as acceptable on at least 7 out of 10 test slides

Summary

Week	Key Deliverable	Owners
1	Bug fixes done, better text/table quality, multi-slide CLI skeleton, judge module started	Rafał, Max, Shaz
2 (<i>Rafał absent</i>)	Image passthrough, improved text rendering, block diagram support, judge at $\geq 80\%$ accuracy	Max, Harsh, Shaz
3	Mimic mode, end-to-end deck pipeline, judge integrated, professional quality review complete	Rafał, Max, Harsh, Shaz